

# COMPSCI 371D Homework 6

## Part 1: Automatic Differentiation

### Automatic Differentiation Basics

```
In [1]: _names = {}

def reset_auto_diff(name=None):
    global _names
    if name is None:
        _names = {}
    else:
        _names.pop(name, None)
```

```
In [2]: class Node:
    def __init__(self, value, gradient=None, name=None, variables=None):
        self.value = value
        self.gradient = gradient
        self.name = name
        self.variables = variables
```

```
In [3]: class Internal(Node):
    def __init__(self, value, gradient, variables):
        super(Internal, self).__init__(value, gradient=gradient, name=None,
                                         variables=variables)

    def __str__(self):
        grad_string = 'Gradient wrt {}: \n{}'.format(self.variables, self.gradient)
        return '{} \n{}'.format(self.value, grad_string)
```

```
In [4]: class Variable(Node):
    def __init__(self, value, name):
        assert name not in _names, 'Different independent variables must have different name'
        assert isinstance(value, int) or isinstance(value, float), 'Variables must be scalars'
        value = float(value)
        _names[name] = name
        gradient = [1.0]
        super(Variable, self).__init__(value, gradient=gradient, name=name, variables=[name])

    def __str__(self):
        return '{}: \n{}'.format(self.name, self.value)
```

```
In [5]: def print_variables(*args, **keywords):
    try:
        print(keywords['title'])
    except KeyError:
        pass
    print(*args, sep='\n\n', end='\n\n')
```

```
In [6]: def scale_list(sequence, factor):  
        return [factor * item for item in sequence]
```

## A Botched Implementation of times

```
In [7]: def bad_times(a, b):  
        assert isinstance(a, Node) and isinstance(b, Node)  
        value = a.value * b.value  
        a_gradient = scale_list(a.gradient, b.value)  
        b_gradient = scale_list(b.gradient, a.value)  
        # The following two + are list concatenation operators  
        gradient = a_gradient + b_gradient  
        variables = a.variables + b.variables  
        return Internal(value, gradient, variables)
```

```
In [8]: reset_auto_diff()  
  
u = Variable(2.0, 'u')  
v = Variable(4.0, 'v')  
print_variables(u, v)  
  
p = bad_times(u, v)  
print_variables(p, title='product')  
  
s = bad_times(u, u)  
print_variables(s, title='square')  
  
u:  
2.0  
  
v:  
4.0  
  
product  
8.0  
Gradient wrt ['u', 'v']:  
[4.0, 2.0]  
  
square  
4.0  
Gradient wrt ['u', 'u']:  
[2.0, 2.0]
```

## Problem 1.1

## Solution

```
In [9]: def merge(*partials):
    derivatives = {}
    for p in partials:
        values = p[0]
        keys = p[1]
        for i in range(len(keys)):
            v = values[i]
            k = keys[i]
            if k in derivatives.keys():
                v = v + derivatives[k]
            derivatives.update({k: v})
    return list(derivatives.values()), list(derivatives.keys())
```

```
In [10]: def times(a, b):
    assert isinstance(a, Node) and isinstance(b, Node)
    value = a.value * b.value
    a_gradient = scale_list(a.gradient, b.value)
    b_gradient = scale_list(b.gradient, a.value)
    gradient, variables = merge((a_gradient, a.variables), (b_gradient, b.variables))
    return Internal(value, gradient, variables)
```

```
In [11]: reset_auto_diff()

u = Variable(2.0, 'u')
v = Variable(4.0, 'v')
print_variables(u, v)

p = times(u, v)
print_variables(p, title='product')

s = times(u, u)
print_variables(s, title='square')
```

u:  
2.0

v:  
4.0

product  
8.0  
Gradient wrt ['u', 'v']:  
[4.0, 2.0]

square  
4.0  
Gradient wrt ['u']:  
[4.0]

## Problem 1.2

## Solution

```
In [12]: def plus(a, b):
          assert isinstance(a, Node) and isinstance(b, Node)
          value = a.value + b.value
          a_gradient = scale_list(a.gradient, 1)
          b_gradient = scale_list(b.gradient, 1)
          gradient, variables = merge((a_gradient, a.variables), (b_gradient, b.variables))
          return Internal(value, gradient, variables)
```

```
In [13]: reset_auto_diff()

u = Variable(2.0, 'u')
v = Variable(4.0, 'v')
print_variables(u, v)

z = plus(u, v)
print_variables(z, title='sum')

twice = plus(u, u)
print_variables(twice, title='twice')
```

u:  
2.0

v:  
4.0

sum  
6.0  
Gradient wrt ['u', 'v']:  
[1.0, 1.0]

twice  
4.0  
Gradient wrt ['u']:  
[2.0]

## Problem 1.3

## Solution

```
In [14]: import numpy as np
          import matplotlib.pyplot as plt
          import math as math
          %matplotlib inline
```

```
In [15]: def cosine(x):
          assert isinstance(x, Node)
          value = math.cos(x.value)
          gradient = scale_list(x.gradient, -math.sin(x.value))
          return Internal(value, gradient, x.variables)
```

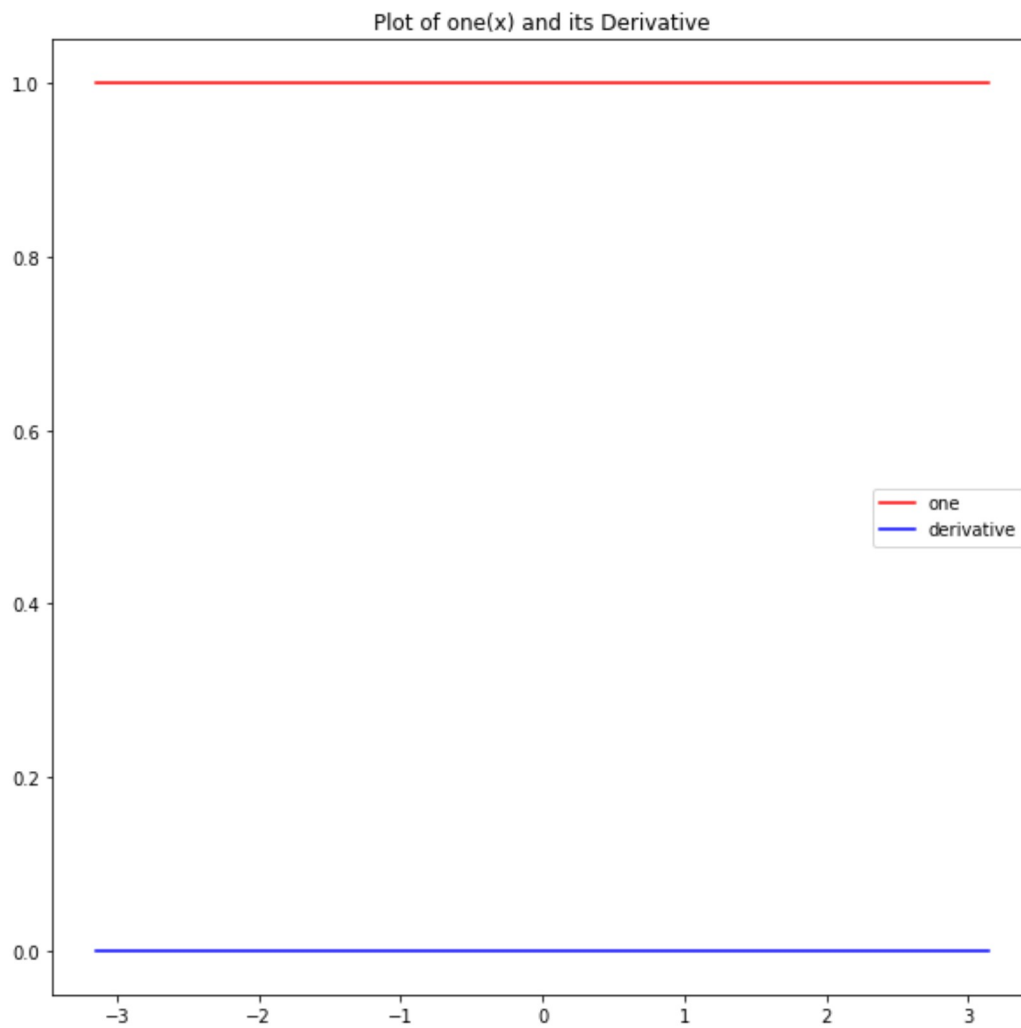
```
In [16]: def sine(x):  
    assert isinstance(x, Node)  
    value = math.sin(x.value)  
    gradient = scale_list(x.gradient, math.cos(x.value))  
    return Internal(value, gradient, x.variables)
```

```
In [17]: def one(x):  
    assert isinstance(x, Node)  
    return plus(times(cosine(x), cosine(x)), times(sine(x), sine(x)))
```

```
In [18]: def derivative(a, name):  
    assert isinstance(a, Node)  
    if name not in a.variables:  
        return 0  
    index = a.variables.index(name)  
    return a.gradient[index]
```

```
In [19]: reset_auto_diff()  
domain = np.linspace(-math.pi, math.pi, endpoint = True)  
one_range = []  
deriv_range = []  
for d in domain:  
    v = Variable(d, str(d))  
    result = one(v)  
    one_range.append(result.value)  
    deriv_range.append(derivative(result, v.variables))
```

```
In [20]: plt.figure(figsize = (10,10))
plt.title("Plot of one(x) and its Derivative")
plt.plot(domain, one_range, '-r', label = 'one')
plt.plot(domain, deriv_range, '-b', label = 'derivative')
x = plt.legend(loc = 'best')
```



## Problem 1.4

## Solution

```

In [21]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

def plot_gradient(f, trange=(-np.pi, np.pi), samples=31):
    t = np.linspace(trange[0], trange[1], samples)
    x_values, y_values = np.meshgrid(t, t)
    z_values = np.zeros_like(x_values)
    g_x = np.zeros((samples, samples))
    g_y = np.zeros((samples, samples))
    reset_auto_diff(name='pg_x')
    reset_auto_diff(name='pg_y')
    x = Variable(0.0, 'pg_x')
    y = Variable(0.0, 'pg_y')
    for i in range(samples):
        for j in range(samples):
            x.value, y.value = x_values[i, j], y_values[i, j]
            z = f(x, y)
            z_values[i, j] = z.value
            g_x[i, j], g_y[i, j] = derivative(z, 'pg_x'), derivative(z, 'pg_y')

    plt.figure(figsize=(8, 8))
    plt.contour(t, t, z_values)
    plt.quiver(t, t, g_x, g_y)
    plt.show()

```

```

In [22]: def trig(x, y):
    return times(times(cosine(x),sine(y)), plus(cosine(x),cosine(times(two, y))))

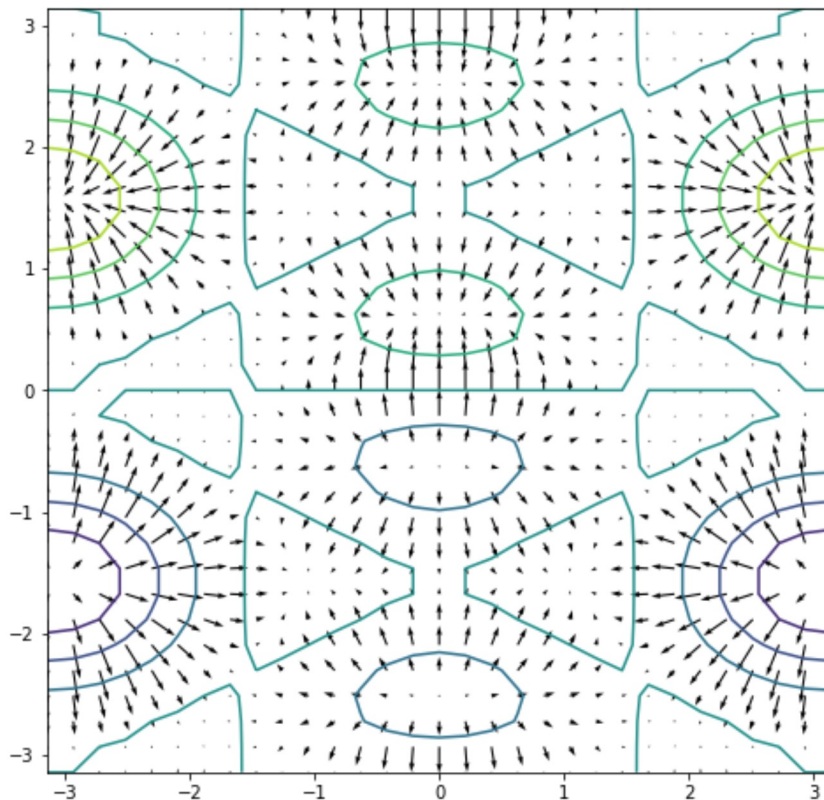
```

```

In [23]: reset_auto_diff()
two = Variable(2.0, 'two')

plot_gradient(trig)

```



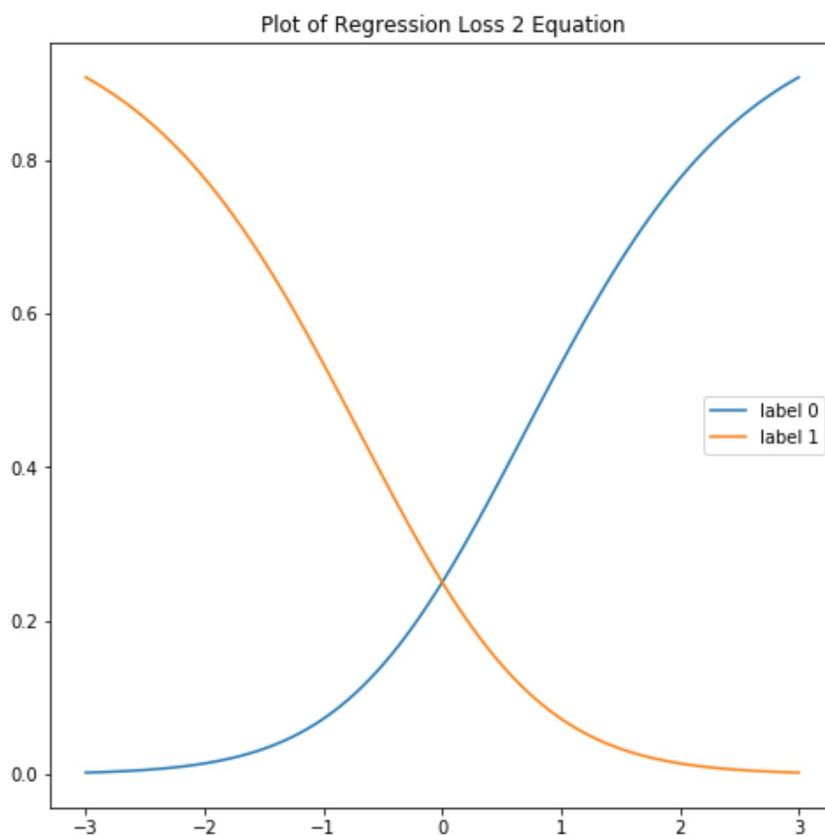
## Part 2: Loss and Convexity

### Problem 2.1

#### Solution

```
In [24]: def regression_loss_2(y, alpha):  
         score = 1 / (1 + np.exp(-alpha))  
         return np.power(score - y, 2)
```

```
In [25]: alphas = np.linspace(-3, 3, 101)  
         plt.figure(figsize=(8, 8))  
         for label in (0, 1):  
             r = regression_loss_2(label, alphas)  
             plt.plot(alphas, r, label='label {}'.format(label))  
         plt.legend(loc='best')  
         t = plt.title('Plot of Regression Loss 2 Equation')
```



### Problem 2.2 (Exam-Style)

#### Solution



By substituting in 0 and  $\alpha$  into our  $r_2$  equation, we get  $\rho(\alpha) = r_2(0, \alpha) = \ell_2(0, s(\alpha)) = p^2$ .

Then  $p^2 = ((1 + e^{-\alpha})^{-1})^2 = (1 + e^{-\alpha})^{-2}$ .

$\rho(\alpha)$  is continuous and differentiable everywhere, so we can find the first derivative:

$$\rho'(\alpha) = \frac{2e^{-\alpha}}{(1 + e^{-\alpha})^3}$$

This is also continuous and differentiable everywhere, so we can find the second derivative:

$$\rho''(\alpha) = \frac{2e^{-2\alpha}(-e^{\alpha} + 2)}{(1 + e^{-\alpha})^4}$$

This second derivative is not positive for all of  $\alpha$  because at  $\alpha \geq \ln(2)$ , the  $e^{\alpha}$  becomes  $\geq 2$  which makes the function negative, so  $\rho(\alpha)$  is not convex.

## Part 3: Logistic Regression and Regularization

```
In [26]: import autograd.numpy as np
         from autograd import grad, jacobian
```

```
In [27]: def logistic(x, v):
         alpha = v[0] + v[1:] * x
         return 1.0 / (1.0 + np.exp(-alpha))
```

```
In [28]: import pickle
         with open('data.pickle', 'rb') as file:
             T = pickle.load(file)
```

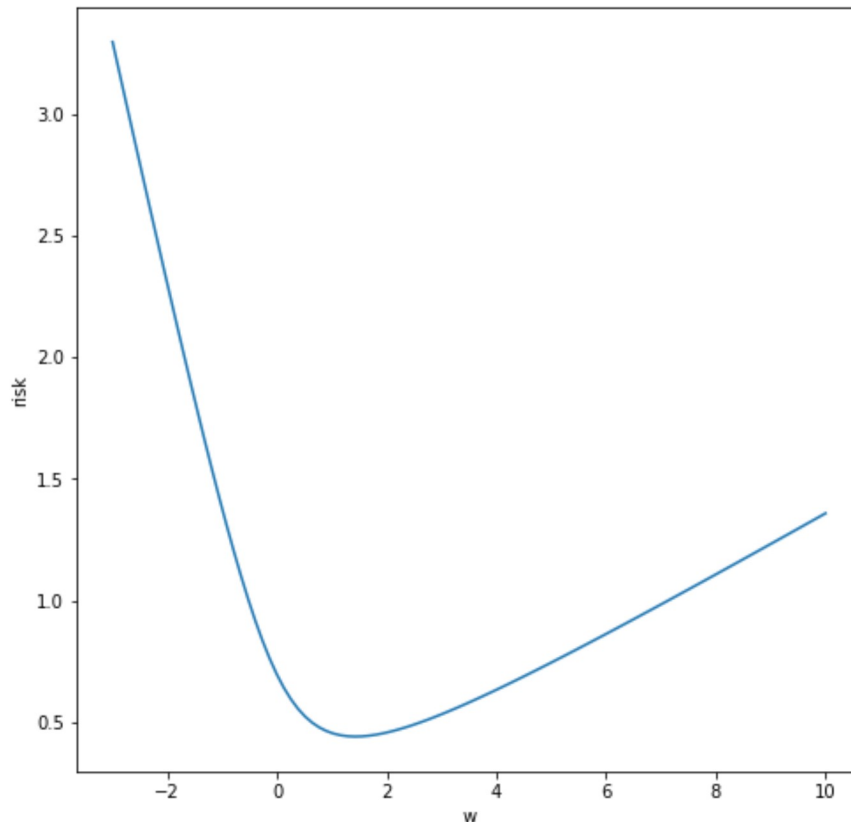
### Problem 3.1

### Solution

```
In [29]: def risk(v, T):
         x = T['x']
         p = []
         small = 1.e-8
         for point in x:
             p.append(logistic(point, v))
         y = T['y']
         risk = []
         for w in range(len(v[1:])):
             total_loss = 0
             for i in range(len(y)):
                 score = np.minimum(1.0 - small, np.maximum(small, p[i][w]))
                 loss = - y[i] * np.log(score) - (1 - y[i]) * np.log(1 - score)
                 total_loss += loss
             risk.append(total_loss/len(y))
         return np.array(risk)
```

```
In [30]: def plot_risks(T, risk=risk):
          bw = []
          bw.append(0)
          for w in np.linspace(-3,10,101):
              bw.append(w)
          v = np.array(bw)
          j = risk(v, T)
          fig = plt.figure(figsize = (8,8))
          plt.plot(bw[1:], j)
          plt.xlabel('w')
          plt.ylabel('risk')
```

```
In [31]: plot_risks(T)
```



## Problem 3.2

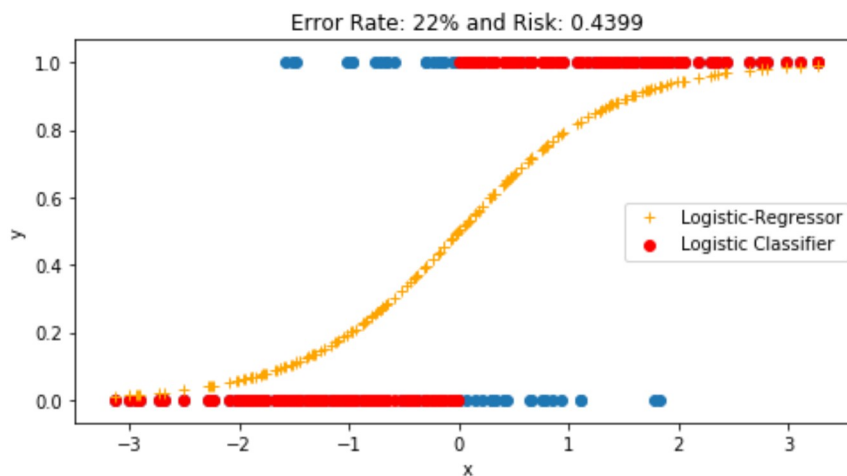
## Solution

```
In [32]: def classifier(x, v):
          y = []
          for data in x:
              if logistic(data, v) > (1/2):
                  y.append(1)
              else:
                  y.append(0)
          return np.array(y)
```

```
In [33]: def performance(v, T):
x = T['x']
h = classifier(x, v)
y = T['y']
loss = 0
for i in range(len(y)):
    if y[i] != h[i]:
        loss+=1
error = loss/len(y) * 100
return error, risk(v, T)
```

```
In [34]: def plot_logistic_regressor(v, T):
fig = plt.figure(figsize = (8,4))
x = T['x']
y = T['y']
plt.scatter(x, y)
regressor = logistic(x, v)
plt.plot(x, regressor, '+', color='orange', label='Logistic-Regressor')
log_classifier = classifier(x, v)
plt.scatter(x, log_classifier, color='red', label='Logistic Classifier')
error, risk = performance(v, T)
plt.xlabel('x')
plt.ylabel('y')
plt.title('Error Rate: %.0f%% and Risk: %.4f' %(error, risk[0]))
plt.legend(loc='best')
```

```
In [35]: v = np.array((0,1.4))
plot_logistic_regressor(v, T)
```



### Problem 3.3

### Solution

```
In [36]: from scipy.optimize import minimize

def learn_logistic_regressor(T, lambda_reg=0.0):
    risk_T = lambda v: risk(v, T) + lambda_reg * np.inner(v, v)
    gradient = grad(risk_T)
    hessian = jacobian(gradient)
    v_0 = np.array((0,0))
    result = minimize(risk_T, v_0, method='Newton-CG', jac=gradient, hess=hessian)
    if result.success:
        v = result.x
        print("b = %.4f \nw = %.4f \nlambda_reg = %.f \nIterations = %.f"
              %(v[0], v[1], lambda_reg, result.nit))
        return result
    else:
        print(result.message)
        return None
```

```
In [37]: v_star = learn_logistic_regressor(T)
```

```
b = -0.0095
w = 1.4299
lambda_reg = 0
Iterations = 6
```

## Problem 3.4

## Solution

```
In [38]: def spread_data(T, spread=0.0):
    x, y = T['x'].copy(), T['y'].copy()
    T_s = {}
    x_s = []
    for i in range(len(y)):
        adjust = -1
        if y[i] == 1:
            adjust = 1
        x_s.append(x[i] + adjust*spread)
    T_s['y'] = y
    T_s['x'] = x_s
    return T_s
```

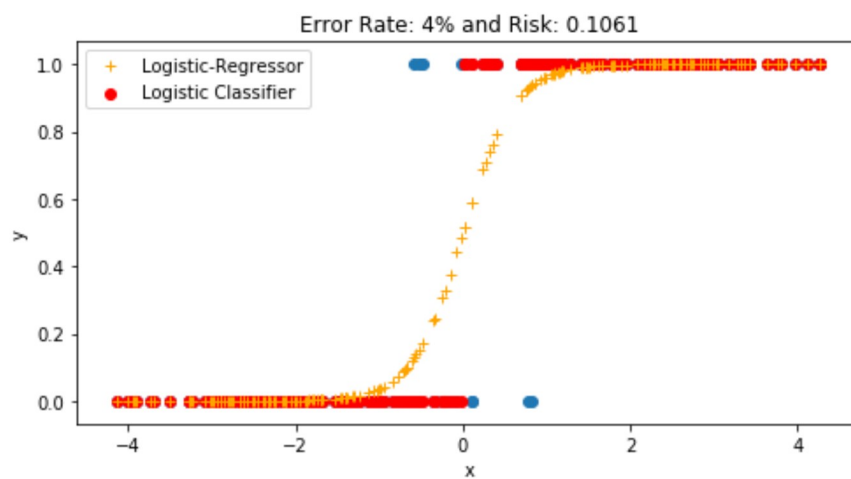
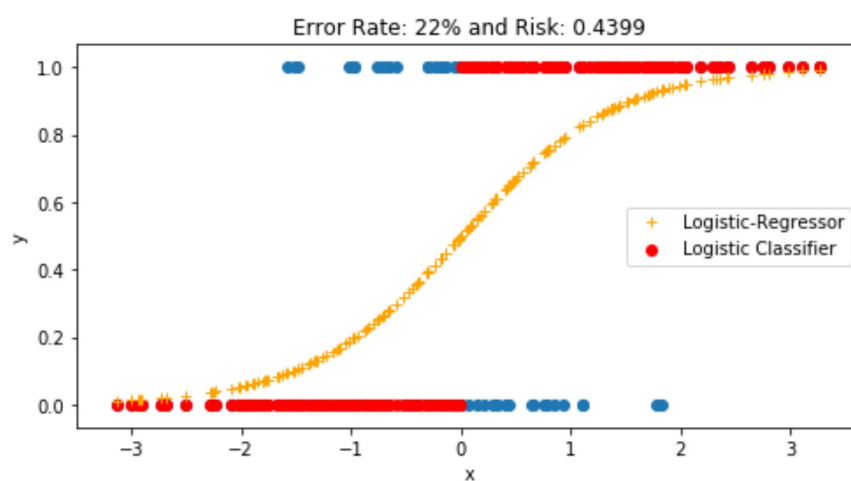
```
In [39]: def spread_experiment(lambda_reg):
    for spread in [0.0, 1.0, 10.0]:
        T_s = spread_data(T, spread=spread)
        result = learn_logistic_regressor(T_s)
        print("Spread Value = %.f" %spread)
        print("-----")
        plot_logistic_regressor(result.x, T_s)
```

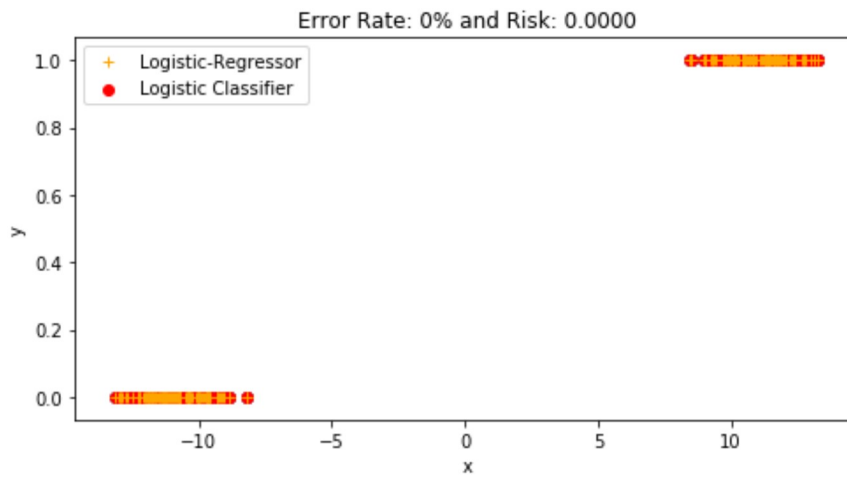
```
In [40]: spread_experiment(0.0)
```

```
b = -0.0095
w = 1.4299
lambda_reg = 0
Iterations = 6
Spread Value = 0
```

```
-----
b = -0.0070
w = 3.2812
lambda_reg = 0
Iterations = 9
Spread Value = 1
```

```
-----
b = -0.0017
w = 1.4882
lambda_reg = 0
Iterations = 15
Spread Value = 10
```

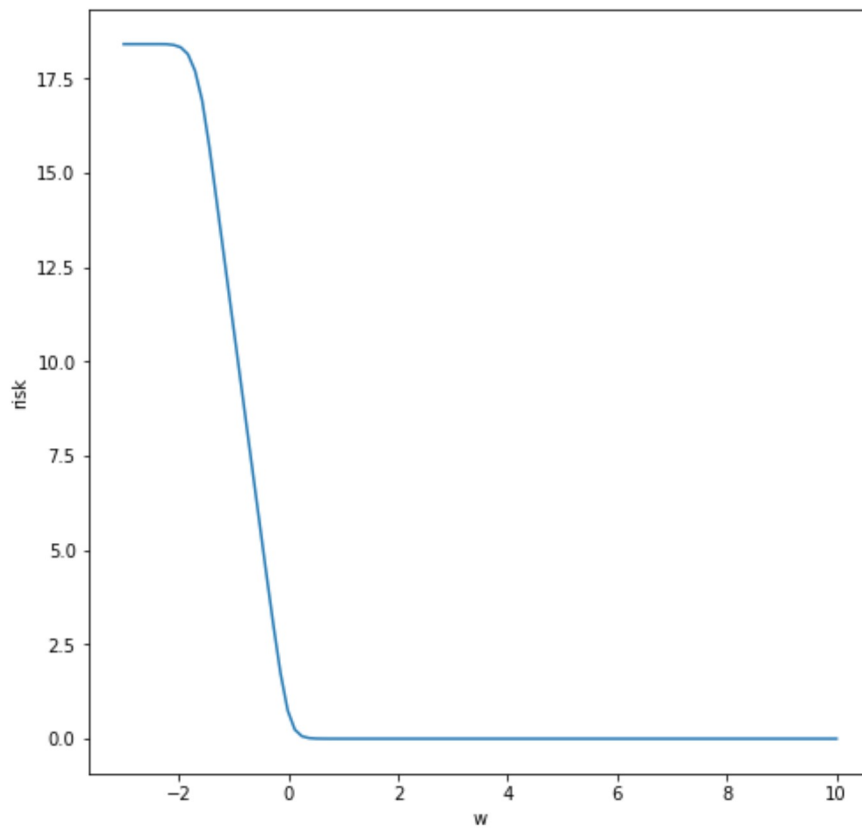




### Problem 3.5

#### Solution

```
In [41]: T_10 = spread_data(T, spread=10.0)
         plot_risks(T_10)
```



### Problem 3.6 (Exam-Style)

## Solution

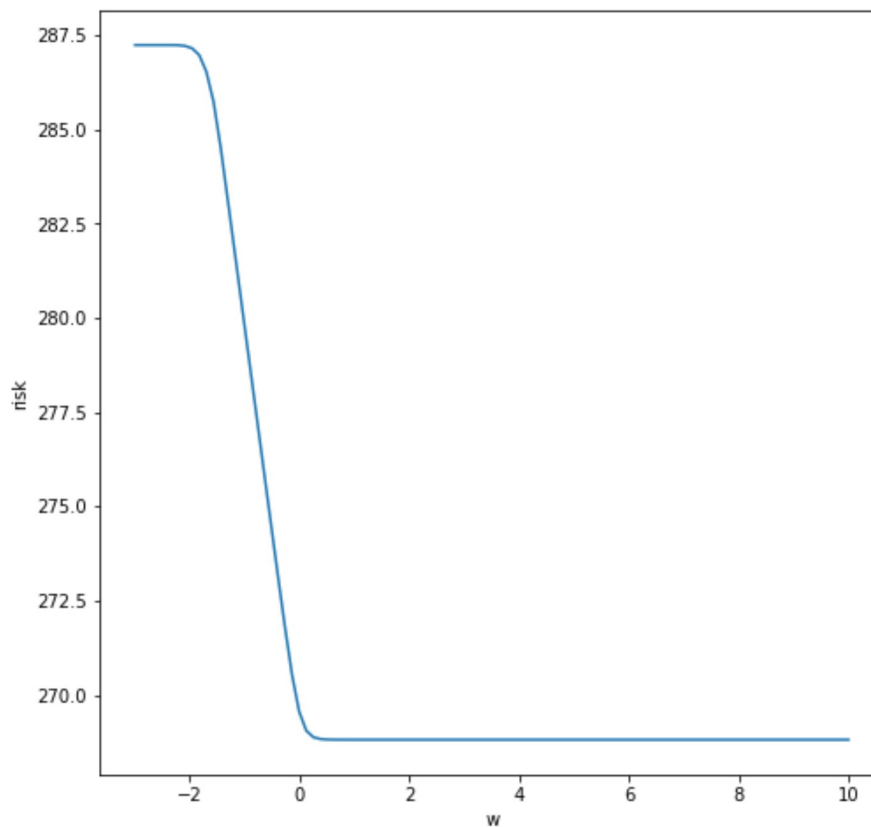
If  $b$  is zero, as  $w$  increases, this will cause the rise of the logistic-regressor function to be sharper. Since the set is separable,  $\hat{w}$  is the value at which the logistic-regression function will perfectly separate the set. So any value of  $w$  beyond  $\hat{w}$  will not reduce the loss, since at  $\hat{w}$  the loss is at 0.

This can be seen in the graph of problem 3.5, at  $w = 1.48$ , the risk is 0, and stays zero after that.

## Problem 3.7

## Solution

```
In [42]: risk_lambda = lambda v, T: risk(v, T) + 0.1 * np.inner(v, v)
         plot_risks(T_10, risk=risk_lambda)
```



The regularization term causes the risk to significantly increase the risk of the function. Before, at  $w = \hat{w}$ , the risk was 0, but now its 270.